

Ablation Study: CUDA Optimization Levels (N=5000)

FP32/FP64 Mixed
Precision ~2.5x

